

Analyze_ab_test_results

January 17, 2019

0.1 Analyze A/B Test Results

0.2 Table of Contents

- Section ??
- Section ??
- Section ??
- Section ??

Introduction

For this project, I will be working to understand the results of an A/B test run by an e-commerce website. My goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

Part I - Probability

To get started, let's import required libraries.

```
In [1]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
import statsmodels.api as sm
%matplotlib inline
#I'm setting the seed to assure you get the same answers on quizzes as I set up
random.seed(42)
```

```
C:\Users\Simin\Anaconda3\lib\site-packages\statsmodels\compat\pandas.py:56: FutureWarning: The
from pandas.core import datetools
```

```
In [2]: # First read in the `ab_data.csv` data and store it in `df`.
df= pd.read_csv('ab_data.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

```

In [4]: # Find the number of rows in the dataset
        df.shape

Out[4]: (294478, 5)

In [5]: # Check to see if there is any duplicate in the dataset
        sum(df.duplicated())

Out[5]: 0

In [6]: # The number of unique users in the dataset
        df['user_id'].nunique()

Out[6]: 290584

In [7]: # The proportion of users converted
        df['converted'].mean()

Out[7]: 0.11965919355605512

In [8]: # The number of times the new_page and treatment don't line up
        df.query(('group == "treatment" and landing_page != "new_page" or group != "treatment"

Out[8]: 3893

In [9]: #Check if the rows have missing values?
        df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
user_id      294478 non-null int64
timestamp    294478 non-null object
group        294478 non-null object
landing_page  294478 non-null object
converted     294478 non-null int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB

```

For the rows where **treatment** is not aligned with **new_page** or **control** is not aligned with **old_page**, I cannot be sure if this row truly received the new or old page. I drop these rows and store the results in df2.

```

In [10]: df2= df.query('group == "treatment" and landing_page == "new_page" or group == "control"

In [11]: # Double Check all of the incorrect rows were removed - this should be 0
        df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].count()

Out[11]: 0

```

```
In [12]: # Number of unique user_id's in df2
df2['user_id'].nunique()
```

```
Out[12]: 290584
```

```
In [13]: #repeated user_id in df2
df2[df2['user_id'].duplicated()]
```

```
Out[13]:
```

	user_id	timestamp	group	landing_page	converted
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

```
In [14]: #the row information for the repeated user_id
df2[df2['user_id']==773192]
```

```
Out[14]:
```

	user_id	timestamp	group	landing_page	converted
1899	773192	2017-01-09 05:37:58.781806	treatment	new_page	0
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

```
In [15]: # Drop the duplicated user
df2.drop([1899,], axis='rows', inplace = True)
```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
In [16]: #Check the results
df2[df2['user_id']==773192]
```

```
Out[16]:
```

	user_id	timestamp	group	landing_page	converted
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

```
In [17]: # The probability of an individual converting regardless of the page they receive?
df2['converted'].mean()
```

```
Out[17]: 0.11959708724499628
```

```
In [18]: # Given that an individual was in the `control` group, what is the probability they c
df2.query('group == "control"')['converted'].mean()
```

```
Out[18]: 0.1203863045004612
```

```
In [19]: # Given that an individual was in the treatment group, what is the probability they c
df2.query('group == "treatment"')['converted'].mean()
```

```
Out[19]: 0.11880806551510564
```

```
In [20]: #The probability that an individual received the new page
(df2.group == "treatment").mean()
```

Out [20]: 0.50006194422266881

Results: Based on the results of what I've done so far, there is no evidence that one page leads to more conversions. Half of the users are in the "treatment" and the other half are in the "control" group, and the rate of conversion is 0.12 for both.

Part II - A/B Test

Because of the time stamp associated with each event, I could technically run a hypothesis test continuously as each observation was observed.

However the hard question is do I stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do I run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, I want to make the decision just based on all the data provided. I assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%. The null and alternative hypotheses are as followed:

p_{old} and p_{new} are the converted rates for the old and new pages.

**** Hypothesis ****

$$H_0 : p_{new} - p_{old} \leq 0$$

$$H_1 : p_{new} - p_{old} > 0$$

2. I'll Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, I'll assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

I'll use a sample size for each page equal to the ones in **ab_data.csv**.

I'll perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

I'll use the cells below to provide the necessary parts of this simulation.

a. The **convert rate** for p_{new} under the null

```
In [21]: p_new = df2['converted'].mean()
         p_new
```

Out [21]: 0.11959708724499628

b. The **convert rate** for p_{old} under the null

```
In [22]: p_old = df2['converted'].mean()
         p_old
```

Out [22]: 0.11959708724499628

c. n_{new}

```
In [23]: n_new = df2.query('landing_page == "new_page")['landing_page'].count()
         n_new
```

Out [23]: 145310

d. n_{old}

```
In [24]: n_old = df2.query('landing_page == "old_page"')['landing_page'].count()
         n_old
```

```
Out [24]: 145274
```

e. Simulate n_{new} transactions with a convert rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

```
In [25]: new_page_converted = np.random.binomial(n=1, p=p_new, size=n_new)
```

f. Simulate n_{old} transactions with a convert rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

```
In [26]: old_page_converted = np.random.binomial(n=1, p=p_old, size=n_old)
```

g. Find $p_{new} - p_{old}$

```
In [27]: new_page_converted.mean() - old_page_converted.mean()
```


```
Out [27]: 0.00051401567146014404
```

h. Simulate 10,000 $p_{new} - p_{old}$ values using the same process in parts a. through g. above and store all 10,000 values in **p_diffs**.

```
In [28]: p_diffs = []
         for _ in range(10000):
             new_bootsamp = np.random.binomial(n=1, p=p_new, size=n_new)
             old_bootsamp = np.random.binomial(n=1, p=p_old, size=n_old)
             p_diffs.append(new_bootsamp.mean() - old_bootsamp.mean())
```

i. Plot a histogram of the **p_diffs**.

```
In [29]: plt.hist(p_diffs);
```



- j. Find the proportion of the **p_diffs** that are greater than the actual difference observed in **ab_data.csv**

```
In [30]: obs_diff = df2.query('group == "treatment"')['converted'].mean()-df2.query('group == "control"')['converted'].mean()
```

```
In [31]: (p_diffs > obs_diff).mean()
```

```
Out[31]: 0.9113999999999999
```

- k. What I computed in part j. was the **p-value**. p_value is the proportion of values in the null distribution which are greater than my observed difference. Since the p-value is a large number, I fail to reject null hypothesis. It means there is no significant difference between the new and old pages.
- l. I could also use a built-in to achieve similar results by calculating the number of conversions for each page, as well as the number of individuals who received each page. n_old and n_new refer the the number of rows associated with the old page and new pages, respectively.

```
In [32]: import statsmodels.api as sm
```

```
convert_old = df2.query('group == "control"')['converted'].sum()
convert_new = df2.query('group == "treatment"')['converted'].sum()
n_old = df2.query('landing_page == "old_page"')['landing_page'].count()
n_new = df2.query('landing_page == "new_page"')['landing_page'].count()
convert_old, convert_new, n_old ,n_new
```

```
Out[32]: (17489, 17264, 145274, 145310)
```

m. Now I use `stats.proportions_ztest` to compute my test statistic and p-value.

```
In [33]: z_score, p_value = sm.stats.proportions_ztest([convert_new,convert_old], [n_new, n_old],
p_value, z_score
```

```
Out[33]: (0.18988337448195103, -1.3109241984234394)
```

I calculated `p_value` under two_tail test, because I left the 'alternative' with default value of "two_sided". `p_value` is different than the one I calculated in j. section. The `p_value` calculated in section j. is for a one_tail test(`p_new - p_old <= 0`)

```
In [34]: import statsmodels.api as sm
z_score, p_value = sm.stats.proportions_ztest([convert_new,convert_old], [n_new, n_old],
p_value, z_score
```

```
Out[34]: (0.90505831275902449, -1.3109241984234394)
```

I set the 'alternative' to "larger", and get the same `p_value` calculated in j. section. with this `p_value` I fail to reject null hypothesis.

Part III - A regression approach

In this final part, will see that the result I achieved in the previous A/B test can also be achieved by performing regression.

- Since each row is either a conversion or no conversion, I should perform Logistic Regression in this case.
- The goal is to use **statsmodels** to fit the regression model I specified in part a. to see if there is a significant difference in conversion based on which page a customer receives. However I first need to create a column for the **intercept**, as well as an **ab_page** column, which is 1 when an individual receives the **treatment** and 0 if **control**.

```
In [35]: df2.head()
```

```
Out[35]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

```
In [36]: df2['intercept'] = 1
df2['ab_page'] = pd.get_dummies(df2['landing_page'])['new_page']
df2.head()
```

```
/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
"""Entry point for launching an IPython kernel.

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
Out [36]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	intercept	ab_page
0	1	0
1	1	0
2	1	1
3	1	1
4	1	0

c. Use **statsmodels** to import logistic regression model.

```
In [37]: log_mod = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])
         results = log_mod.fit()
```

Optimization terminated successfully.
Current function value: 0.366118
Iterations 6

d. Provide the summary of the results.

```
In [38]: results.summary()
```

```
Out [38]: <class 'statsmodels.iolib.summary.Summary'>
        """
```

Logit Regression Results					
=====					
Dep. Variable:	converted	No. Observations:	290584		
Model:	Logit	Df Residuals:	290582		
Method:	MLE	Df Model:	1		
Date:	Wed, 30 May 2018	Pseudo R-squ.:	8.077e-06		
Time:	17:38:30	Log-Likelihood:	-1.0639e+05		
converged:	True	LL-Null:	-1.0639e+05		
		LLR p-value:	0.1899		
=====					
	coef	std err	z	P> z	[0.025 0.975]


```
-----
intercept      -1.9888      0.008   -246.669      0.000      -2.005      -1.973
ab_page        -0.0150      0.011    -1.311      0.190      -0.037      0.007
=====
"""
```

The p-value associated with ab_page is 0.19.

In linear regression, I discuss if there is a significant linear relationship between the dependent variable/column X(in my case ab_page) and independent variable/column Y(converted). The null hypothesis associated with regression is that there is no significant relationship between ab_page and converted columns which means conversion rate is not dependent on the page visited by user(old/new page). and the alternative hypothesis is that there is a relationship between ab_page and converted columns. These cause a two_tail hypothesis.

**** Hypothesis: ****

$$H_0 : p_{new} - p_{old} = 0$$

$$H_1 : p_{new} - p_{old} \neq 0$$

The p_value I calculated in the regression model is under a two_tail test. In Part II, section m. under a two_tail test I found the same value for p_value, however under a one_tail hypothesis the p_value was 0.91.

**** f. Is it a good idea to consider other factors to add into my regression model? what are the pros and cons? **** Considering other factors to add into my regression model could help us to predict the result better or finding more associated variables, on the other hand it could cause complexity in interpreting the results or collinearity which is correlation between predictor variables that can lead to coefficients being flipped from the direction I expect from simple linear regression.

g. Now along with testing if the conversion rate changes for different pages, I also add an effect based on which country a user lives. I will read in the **countries.csv** dataset and merge together my datasets on the appropriate rows.

I want to see if country had an impact on conversion?

```
In [39]: #reading countries.csv file
countries_df = pd.read_csv('countries.csv')
countries_df.head()
```

```
Out[39]:   user_id country
0    834778      UK
1    928468      US
2    822059      UK
3    711597      UK
4    710616      UK
```

```
In [40]: #checking for null values in countries.df dataset
countries_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 290584 entries, 0 to 290583
```

```
Data columns (total 2 columns):
user_id      290584 non-null int64
country      290584 non-null object
dtypes: int64(1), object(1)
memory usage: 4.4+ MB
```

```
In [41]: #checking if I have the same row number in df2 as countries_df
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 290584 entries, 0 to 294477
Data columns (total 7 columns):
user_id      290584 non-null int64
timestamp    290584 non-null object
group        290584 non-null object
landing_page 290584 non-null object
converted     290584 non-null int64
intercept    290584 non-null int64
ab_page      290584 non-null uint8
dtypes: int64(3), object(3), uint8(1)
memory usage: 15.8+ MB
```

```
In [42]: #Joining the two datasets and saving the results in df3
df3 = df2.set_index('user_id').join(countries_df.set_index('user_id'))
df3.head()
```

```
Out [42]:
```

		timestamp	group	landing_page	converted	\
user_id						
851104	2017-01-21	22:11:48.556739	control	old_page	0	
804228	2017-01-12	08:01:45.159739	control	old_page	0	
661590	2017-01-11	16:55:06.154213	treatment	new_page	0	
853541	2017-01-08	18:28:03.143765	treatment	new_page	0	
864975	2017-01-21	01:52:26.210827	control	old_page	1	

	intercept	ab_page	country
user_id			
851104	1	0	US
804228	1	0	US
661590	1	1	US
853541	1	1	US
864975	1	0	US

```
In [43]: #finding the existing country in my dataset to help us create necessary dummy variables
df3.country.value_counts()
```

```
Out [43]: US      203619
          UK       72466
```

```
CA      14499
Name: country, dtype: int64
```

```
In [44]: #creating dummy variables for exising countries
df3[['CA', 'UK', 'US']] = pd.get_dummies(df3['country'])
df3 = df3.drop('CA', axis=1)
df3.head()
```

```
Out [44]:
```

	timestamp	group	landing_page	converted	\
user_id					
851104	2017-01-21 22:11:48.556739	control	old_page	0	
804228	2017-01-12 08:01:45.159739	control	old_page	0	
661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
864975	2017-01-21 01:52:26.210827	control	old_page	1	

	intercept	ab_page	country	UK	US
user_id					
851104	1	0	US	0	1
804228	1	0	US	0	1
661590	1	1	US	0	1
853541	1	1	US	0	1
864975	1	0	US	0	1

h. Though I have now looked at the individual factors of country and page on conversion, I would now like to look at an interaction between page and country to see if there is significant effects on conversion.

```
In [45]: log_mod = sm.Logit(df3['converted'], df3[['intercept', 'ab_page', 'UK', 'US']])
results = log_mod.fit()
results.summary()
```

```
Optimization terminated successfully.
Current function value: 0.366113
Iterations 6
```

```
Out [45]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                Logit Regression Results
=====
Dep. Variable:                  converted    No. Observations:                  290584
Model:                            Logit      Df Residuals:                      290580
Method:                           MLE        Df Model:                          3
Date:                            Wed, 30 May 2018    Pseudo R-squ.:                    2.323e-05
Time:                            17:38:32      Log-Likelihood:                   -1.0639e+05
converged:                        True          LL-Null:                         -1.0639e+05
                                      LLR p-value:                        0.1760
=====
```

	coef	std err	z	P> z	[0.025	0.975]
intercept	-2.0300	0.027	-76.249	0.000	-2.082	-1.978
ab_page	-0.0149	0.011	-1.307	0.191	-0.037	0.007
UK	0.0506	0.028	1.784	0.074	-0.005	0.106
US	0.0408	0.027	1.516	0.130	-0.012	0.093

=====

"""

According to the results, because all the p-values are greater than 0.05, there isn't any statistically significant relationship between the variables and the response variable(converted)

```
In [46]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Analyze_ab_test_results_notebook.ipynb'])
```

```
Out[46]: 0
```